

Task: Execute WordMedian , WordMean , WordStandardDeviation programs using hadoop-mapreduce-examples-2.9.0.jar file present in your AcadGild VM. Refer path below.
/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce

```
[acadgild@localhost Desktop]$ hadoop fs -cat /Customer_Input.txt
18/09/04 17:28:49 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
1234      Fname      LName      India
1234      Fname      LName      India
1234      Fname      LName      India
1234      Fname      LName      India
[acadgild@localhost Desktop]$
```

WordMedian:

```

acadgild@localhost:~/Desktop
File Edit View Search Terminal Help
[acadgild@localhost Desktop]$ hadoop jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordmedian /Customer_Input.txt /AssignmentOutputYarn/
18/09/04 17:10:12 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/09/04 17:10:14 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
org.apache.hadoop.mapred.FileAlreadyExistsException: Output directory hdfs://localhost:8020/AssignmentOutput already exists
    at org.apache.hadoop.mapreduce.lib.output.FileOutputFormat.checkOutputSpecs(FileOutputFormat.java:146)
    at org.apache.hadoop.mapreduce.JobSubmitter.checkSpecs(JobSubmitter.java:267)
    at org.apache.hadoop.mapreduce.JobSubmitter.submitJobInternal(JobSubmitter.java:140)
    at org.apache.hadoop.mapreduce.Job$10.run(Job.java:1297)
    at org.apache.hadoop.mapreduce.Job$10.run(Job.java:1294)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at org.apache.hadoop.mapreduce.Job.submit(Job.java:1294)
    at org.apache.hadoop.mapreduce.Job.waitForCompletion(Job.java:1315)
    at org.apache.hadoop.examples.WordMedian.run(WordMedian.java:194)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:70)
    at org.apache.hadoop.examples.WordMedian.main(WordMedian.java:171)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:498)
    at org.apache.hadoop.util.ProgramDriver$ProgramDescription.invoke(ProgramDriver.java:71)
    at org.apache.hadoop.util.ProgramDriver.run(ProgramDriver.java:144)
    at org.apache.hadoop.examples.ExampleDriver.main(ExampleDriver.java:74)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:498)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:221)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
[acadgild@localhost Desktop]$ hadoop jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordmedian /Customer_Input.txt /AssignmentOutputYarn/
[acadgild@localhost Desktop]$ hadoop jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordmedian /Customer_Input.txt /AssignmentOutputYarn/
18/09/04 17:10:38 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/09/04 17:10:39 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032

```

```

Virtual memory (bytes) snapshot=4118
Total committed heap usage (bytes)=1
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=92
File Output Format Counters
Bytes Written=9
The median is: 5
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost Desktop]$

```

```

[acadgild@localhost Desktop]$ hadoop fs -ls /AssignmentOutputYarn/
18/09/04 17:27:57 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 acadgild supergroup 0 2018-09-04 17:11 /AssignmentOutputYarn/_SUCCESS
-rw-r--r-- 1 acadgild supergroup 9 2018-09-04 17:11 /AssignmentOutputYarn/part-r-00000
You have new mail in /var/spool/mail/acadgild

```

WordMean:

```
acadgild@localhost:~/Desktop
File Edit View Search Terminal Help
[acadgild@localhost Desktop]$ hadoop jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordmean /Customer_Input.txt /AssignmentOutputYarn1/
18/09/04 17:16:59 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/09/04 17:17:00 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/09/04 17:17:02 INFO input.FileInputFormat: Total input paths to process : 1
18/09/04 17:17:02 INFO mapreduce.JobSubmitter: number of splits:1
18/09/04 17:17:03 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1536061112797_0002
18/09/04 17:17:03 INFO impl.YarnClientImpl: Submitted application application_1536061112797_0002
18/09/04 17:17:03 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1536061112797_0002/
18/09/04 17:17:03 INFO mapreduce.Job: Running job: job_1536061112797_0002
18/09/04 17:17:13 INFO mapreduce.Job: Job job_1536061112797_0002 running in uber mode : false
18/09/04 17:17:13 INFO mapreduce.Job: map 0% reduce 0%
18/09/04 17:17:20 INFO mapreduce.Job: map 100% reduce 0%
18/09/04 17:17:28 INFO mapreduce.Job: map 100% reduce 100%
18/09/04 17:17:28 INFO mapreduce.Job: Job job_1536061112797_0002 completed successfully
18/09/04 17:17:29 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=39
    FILE: Number of bytes written=215429
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=197
    HDFS: Number of bytes written=19
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=4692
    Total time spent by all reduces in occupied slots (ms)=4947
    Total time spent by all map tasks (ms)=4692
    Total time spent by all reduce tasks (ms)=4947
    Total vcore-milliseconds taken by all map tasks=4692
    Total vcore-milliseconds taken by all reduce tasks=4947
    Total megabyte-milliseconds taken by all map tasks=4804608
```

```
Reduce input groups=2
Reduce shuffle bytes=39
Reduce input records=2
Reduce output records=2
Spilled Records=4
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=155
CPU time spent (ms)=1310
Physical memory (bytes) snapshot=297955328
Virtual memory (bytes) snapshot=4118224896
Total committed heap usage (bytes)=170004480

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=92
File Output Format Counters
  Bytes Written=19
```

The mean is: 4.75

You have new mail in /var/spool/mail/acadgild
[acadgild@localhost Desktop]\$

```
[acadgild@localhost Desktop]$ hadoop fs -ls /AssignmentOutputYarn1/
18/09/04 17:26:42 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 acadgild supergroup 0 2018-09-04 17:17 /AssignmentOutputYarn1/_SUCCESS
-rw-r--r-- 1 acadgild supergroup 19 2018-09-04 17:17 /AssignmentOutputYarn1/part-r-00000
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost Desktop]$
```

WordStandardDeviation:

```
acadgild@localhost:~/Desktop
File Edit View Search Terminal Help
[acadgild@localhost Desktop]$ hadoop jar /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar wordstandarddeviation /Customer Input.txt /AssignmentOutputYarn2/
18/09/04 17:21:38 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/09/04 17:21:39 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/09/04 17:21:41 INFO input.FileInputFormat: Total input paths to process : 1
18/09/04 17:21:41 INFO mapreduce.JobSubmitter: number of splits:1
18/09/04 17:21:41 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1536061112797_0003
18/09/04 17:21:42 INFO impl.YarnClientImpl: Submitted application application_1536061112797_0003
18/09/04 17:21:42 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1536061112797_0003/
18/09/04 17:21:42 INFO mapreduce.Job: Running job: job_1536061112797_0003
18/09/04 17:21:51 INFO mapreduce.Job: Job job_1536061112797_0003 running in uber mode : false
18/09/04 17:21:51 INFO mapreduce.Job: map 0% reduce 0%
18/09/04 17:21:58 INFO mapreduce.Job: map 100% reduce 0%
18/09/04 17:22:07 INFO mapreduce.Job: map 100% reduce 100%
18/09/04 17:22:08 INFO mapreduce.Job: Job job_1536061112797_0003 completed successfully
18/09/04 17:22:08 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=56
    FILE: Number of bytes written=215623
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=197
    HDFS: Number of bytes written=30
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=5009
    Total time spent by all reduces in occupied slots (ms)=5444
    Total time spent by all map tasks (ms)=5009
    Total time spent by all reduce tasks (ms)=5444
    Total vcore-milliseconds taken by all map tasks=5009
    Total vcore-milliseconds taken by all reduce tasks=5444
    Total megabyte-milliseconds taken by all map tasks=5129216
```

```
Merged Map outputs=1
GC time elapsed (ms)=174
CPU time spent (ms)=1340
Physical memory (bytes) snapshot=298815488
Virtual memory (bytes) snapshot=4118224896
Total committed heap usage (bytes)=170004480

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=92
File Output Format Counters
  Bytes Written=30
```

The standard deviation is: 0.4330127018922193

You have new mail in /var/spool/mail/acadgild

```
[acadgild@localhost Desktop]$
```

```
[acadgild@localhost Desktop]$ hadoop fs -ls /AssignmentOutputYarn2/
18/09/04 17:25:18 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 acadgild supergroup          0 2018-09-04 17:22 /AssignmentOutputYarn2/_SUCCESS
-rw-r--r-- 1 acadgild supergroup       30 2018-09-04 17:22 /AssignmentOutputYarn2/part-r-00000
[acadgild@localhost Desktop]$
```